

---

# Proof of Latency Using a Verifiable Delay Function

---

Department of Future Technologies

Master's Thesis  
University of Turku  
Department of Future Technologies

2020  
Jani Anttonen

UNIVERSITY OF TURKU  
Department of Information Technology

JANI ANTTONEN: Proof of Latency Using a Verifiable Delay Function

Master's Thesis, 44 p., 0 app. p.  
Department of Future Technologies  
May 2020

---

Tarkempia ohjeita tiivistelmäsivun laadintaan löytyy opiskelijan yleisoppaasta, josta alla lyhyt katkelma.

Bibliografisten tietojen jälkeen kirjoitetaan varsinainen tiivistelmä. Sen on oletettava, että lukijalla on yleiset tiedot aiheesta. Tiivistelmän tulee olla ymmärrettävissä ilman tarvetta perehtyä koko tutkielmaan. Se on kirjoitettava täydellisinä virkkeinä, väliotsakeluettelona. On käytettävä vakiintuneita termejä. Viittauksia ja lainauksia tiivistelmään ei saa sisällyttää, eikä myöskään tietoja tai väitteitä, jotka eivät sisälly itse tutkimukseen. Tiivistelmän on oltava mahdollisimman ytimekäs n. 120 – 250 sanan pituinen itsenäinen kokonaisuus, joka mahtuu ykkäsvälillä kirjoitettuna vaivatta tiivistelmäsivulle. Tiivistelmässä tulisi ilmetä mm. tutkielman aihe tutkimuksen kohde, populaatio, alue ja tarkoitus käytetyt tutkimusmenetelmät (mikäli tutkimus on luonteeltaan teoreettinen ja tiettyyn kirjalliseen materiaaliin, on mainittava tärkeimmät lähdeoteokset; mikäli on luonteeltaan empiirinen, on mainittava käytetyt menetit) keskeiset tutkimustulokset tulosten perusteella tehdyt päätelmät ja toimenpidesuosituksat asiasanat

Keywords: list, of, keywords

# Contents

## List of Figures

## List of Tables

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Cryptography</b>	<b>4</b>
2.1	RSA . . . . .	4
2.2	Elliptic Curve Cryptography . . . . .	5
2.3	Groups of Unknown Order . . . . .	5
2.4	Cryptographic Proofs . . . . .	5
<b>3</b>	<b>Peer-to-Peer Networking</b>	<b>7</b>
3.1	Ad Hoc and Zero Configuration Networking . . . . .	10
3.2	Distributed Hash Tables . . . . .	11
3.2.1	Kademlia . . . . .	11
<b>4</b>	<b>Verifiable Delay Functions</b>	<b>13</b>
4.1	History . . . . .	14
4.2	Use Cases . . . . .	15
4.3	Variations . . . . .	16
4.4	Similar Constructs . . . . .	16

4.4.1	Proto-VDF . . . . .	16
4.4.2	Classic Slow Functions . . . . .	16
4.5	Hardware Developments . . . . .	17
<b>5</b>	<b>Proof of Latency</b>	<b>19</b>
5.1	Use Cases . . . . .	19
5.1.1	Dynamic P2P Routing . . . . .	20
5.1.2	Benchmark Queries . . . . .	20
5.2	Role of Latency in Distributed Systems . . . . .	21
5.3	Network Hops Increase Latency . . . . .	23
5.4	Protocol Description . . . . .	23
5.5	Attack Vectors . . . . .	26
5.5.1	Advertising Dishonest Peers . . . . .	26
5.5.2	Performance Matching . . . . .	27
5.6	Protecting Against Performance Matching . . . . .	27
5.6.1	Hiding Information of VDF Results . . . . .	27
5.6.2	Web of Trust . . . . .	28
5.6.3	Peer Scoring . . . . .	28
5.6.4	Switch Responsibilities . . . . .	28
5.6.5	Results . . . . .	29
<b>6</b>	<b>Proof of Concept</b>	<b>30</b>
6.1	Tests . . . . .	36
6.1.1	Unit Tests . . . . .	36
6.1.2	End to End Tests . . . . .	37
<b>7</b>	<b>Future Considerations</b>	<b>38</b>
<b>8</b>	<b>Conclusion</b>	<b>39</b>



# List of Figures

3.1	Example topography of an unstructured overlay network . . . . .	8
3.2	Example topography of a structured overlay network . . . . .	9
5.1	Protocol Diagram, Proof of Latency . . . . .	24

## **List of Tables**

# Chapter 1

## Introduction

Computer applications, whether they are on the public Internet or on a private network, have long preferred a client-server model of communication. In this model, the client application sends requests to the server, and the server responds. Applications are today, however, increasingly moving towards a distributed, peer-to-peer (P2P) networked model, where every peer, whether it is an entire computer or merely a process running on one, serves equally as the both sides of the client-server model. Many uses of P2P do not even get communicated to the end user. For example, the music subscription service Spotify's protocol has been designed to combine server and P2P streaming to improve scalability by decreasing the load on Spotify's servers and bandwidth resources.[1]

The aspects of of peer-to-peer networking that have recently got attention have been cryptocurrency and blockchain technologies, categorized under the roof term of distributed ledger technology. Peer-to-peer has not been really shown in the public light as anything more than a technology to work around regulation and for doing lawless activities. This is partly because before its use in blockchain applications, it was popularized for its use in file sharing applications like Bittorrent, which it still sees use for.

Blockchain networks need a way of synchronizing the latest state of the blockchain globally between a number of peers on a P2P network. Since the blockchain data model is sequential and all recorded history is immutable, an algorithm is needed to reach total



synchronization between peers. This kind of algorithm is called a consensus algorithm. This problem is not unique to blockchains, it is relevant in distributed databases as well, but public blockchains have raised new issues that have sparked an ongoing development effort for new kinds of consensus algorithms.

Proof of Work is the most used consensus algorithm in public blockchains today, including Bitcoin, of in which whitepaper it was first described in 2008. [2] New algorithms have been introduced since to battle Proof of Work's resource intensiveness, including Proof of Stake, which requires network nodes participating in the voting of new blocks to stake a part of their assets as a pawn. Simply this means handing the control of some of the currency owned to the consensus algorithm if the peer wants to participate in the consensus. If a voter gets labeled as malicious, faulty, or absent by a certain majority, it can get slashed, losing all or a part of the staked asset in the process. This serves as an incentive for honest co-operation, with sufficient computation resources.

One problem with Proof of Stake is that the block generation votes are not done globally, but by a selected group of peers called the validators, which vote for the contents of proposed blocks that are generated by just one peer at a time, selected as the block generator. The validators are usually selected randomly. This has generated an increasing demand for verifiable public randomness that is pre-image resistant, meaning the output of the algorithm generating the randomness cannot be influenced before evaluation by input. This created a need for an algorithm that would prevent multiple malicious actors from being selected to vote at once. A cure for this problem is called a verifiable delay function, a VDF in short.

In 2018, two research papers were released independently with similar formalizations of a VDF. [3, 4] VDF is an algorithm that requires a specified number of sequential steps to evaluate, but produces a proof that can be efficiently and publicly verified. [5] To achieve pre-image resistance, a VDF is sequential in nature, and cannot be sped up by parallel processing. There are multiple formulations of a VDF, and not all even have a generated

proof, instead using parallel processing with graphics processors to check that the delay function has been calculated correctly. [6] This bars less powerful devices, like embedded devices, from verifying the VDF's result efficiently. Thus, generating a proof that requires little time to verify is more ideal. [5]

Using a verifiable delay function, I propose a novel algorithm for producing a publicly verifiable proof of network latency and difference in computation resources between two participants in a peer-to-peer network. This proof can be used for dynamic routing to reduce latency between peers, and for making eclipse attacks<sup>1</sup> harder to achieve.

---

<sup>1</sup>Eclipse attack means polluting the target's routing table restricting the target's access to the rest of the network, which opens up other attack possibilities regarding consensus algorithms.

# Chapter 2

## Cryptography

Modern cryptography is based on relatively few mathematical fundamentals. The most dominant have historically been prime numbers and modular multiplication, but elliptic curve cryptography is increasingly considered to be the most secure method out there to keep things colloquial.

Modular multiplication with large numbers has a useful property that you can exchange encrypted data without the participants knowing each other's private keys, but requiring a computation that is theoretically almost impossible to break. Prime numbers and elliptic curves can be used to create the variables used in modular multiplication, providing the basis for the robustness of the cryptosystem.

### 2.1 RSA

RSA is named after its discoverers – Rivest, Shamir, Adleman. It is an asymmetric public-key cryptosystem, meaning that it is based on a keypair, in which one is a public and the other is a secret. The public key is used to encrypt, and the secret key is used to decrypt. This means that everyone who has the public key can encrypt data and that encrypted data is practically impossible to decrypt without the secret key. This works due to the fact that it is hard to factor the product of two large prime numbers.

RSA is an arithmetic trick that creates a mathematical object called a trapdoor permutation. Trapdoor permutation is a function that transforms a number  $x$  to a number  $y$  in the same range, in a way that computing  $y$  from  $x$  is easy using the public key but computing  $x$  from  $y$  is practically impossible without knowing the private key. The private key is the trapdoor.[7]

Some institutions also have been publishing these products of two large prime numbers, claiming they have discarded the two prime numbers that were used in the creation. These are published as public puzzles with prizes for correct factorizations as high as 200000 US dollars, but the larger ones can also be used in cryptography to remove a trusted setup. These products are called RSA numbers, of which RSA-1024 and RSA-2048 are widely used. They can be considered relatively safe for production use, since the latest broken RSA number at the time of writing is RSA-250. Still, in practice, the keys could be still stored by the number's creator, requiring trust towards the institutions or individuals who have published them.

## **2.2 Elliptic Curve Cryptography**

## **2.3 Groups of Unknown Order**

## **2.4 Cryptographic Proofs**

Cryptographic proofs are proofs that depend on the trapdoor nature of cryptographic functions. The most well-known proofs, which are not necessarily thought of as such are cryptographic signatures. Given a public key, a computer in possession of its corresponding private key can produce a signature of any given data, implying together with the data that the computer has seen and processed the data.

Proofs can be private or public. A cryptographic proof can be categorized as public, if

a verifier can gather all information required to verify the proof from the transcript of the proof itself, and verify the proof to be correct. Now, since classical cryptography is based on the fact that a computation is asymmetric — being harder to compute the other way around, a cryptographic proof is still probabilistic in nature, and the security of it based on the parameters used.

## Chapter 3

# Peer-to-Peer Networking

Peer-to-peer networking, P2P in short, means that two or more devices communicate with each other with both serving as a client and a server simultaneously. P2P networks vary in scope a lot. Some networks are global, while some are as small as a couple of devices in a local area network, which is the case in printer sharing and some IoT installations. P2P networks can be standalone or rely on some existing infrastructure, like the internet. The P2P networks that rely on preexisting infrastructure are called overlay networks.

While overlay networks' addressing and routing is based on TCP/IP underneath, they must introduce a separate routing scheme, like a DHT, that works as a search engine for peers on the network. Since TCP/IP only focuses on addressing the computers and not the applications or resources they offer, it can't function as a peer discovery scheme.

Overlay P2P networks must somehow signal the initial peers, also called bootstrap or introductory peers, to connect to at first. This can be done either by including the bootstrap peers' IP addresses in the application code itself or by providing the same information on a regular web page. The bootstrap peers could be also signaled by word-of-mouth between a group of people, or with wireless broadcast, which I will touch a bit on the next section. This said, usually every peer can function as a bootstrap peer, so knowing the connection info of a single peer on the network can potentially introduce you into the network at large.




Figure 3.1: Example topography of an unstructured overlay network

A simple way to create a P2P network would be to share a single file with anyone who connects to you that has all the IP addresses you've ever connected with. This would be an example of an unstructured P2P network, which could work relatively well in a closed setting, like inside a LAN if there's no support for multicast DNS, or mDNS in short.

In contrast to unstructured networks, where peers are connected to at random, structured networks have some logic that is used to form a structured overlay network rather than just connecting to peers as they come.




Figure 3.2: Example topography of a structured overlay network



### 3.1 Ad Hoc and Zero Configuration Networking

Multicast DNS was proposed by Apple in 2013 [8] as a way of discovering peers in a local area network in a zero-configuration manner. It is used today for resource sharing, such as sharing printers. Multicast DNS does not work outside local area networks, since it works by associating names with IP addresses like regular DNS does. The problem is that these names are not guaranteed to be unique, and therefore can be spoofed. If there are two clients with the same name, the first one to respond with its IP address to a query wins. [9] The security of zero-configuration and ad hoc networking must rely on cryptographic identities, so that a peer can verify itself with public-key cryptography, that makes the peers on the network practically unique and thus hard to spoof.

Zero-configuration networking in an unconstrained, global setting is possible with radios, using either dedicated meshnet radios like GoTenna [10] or Helium [11], or by using the antennas inside mobile devices to form a network. These types of networks are usually called meshnets or ad hoc networks. Basically, Walkie-Talkies are the simplest form of a P2P network. Problems can arise due to geography blocking signals, or when you want to cross large distances with the transmitting capability of a pocket device. Mesh networking has been used most famously in protests worldwide by using a smartphone app called Firechat. [12]

Ad hoc mesh networks have a natural metric for latency, signal strength, and thus will not benefit from dynamic routing offered by DHTs as much. They can rely on Bluetooth RSSI<sup>1</sup> as a metric, or triangulate distances by cooperating with multiple peers. This type of system is used to locate emergency calls and in contact tracing. [13] Mesh networks, while very much peer to peer and not relying on existing infrastructure like overlay networks, still need routing if one wants to reach peers that are not in the operating range of the communication method used.

---

<sup>1</sup>Received Signal Strength Indication

## 3.2 Distributed Hash Tables

Distributed hash tables are a way of pointing content to peers in a distributed network. In addition to indexing content in content-addressed networks like IPFS [14], they can function as routing tables. A hash table is a key-value mapping from a to b. What makes them distributed is the fact that the data stored is meant to be distributed between peers, with not a single peer keeping all the available data in its hash table, but relaying queries for resources it can't reach to other peers on the network. There are multiple versions of DHTs with different methods for prioritizing certain peers, using tree structures, sorting by identifiers, using computational trust et cetera.

Most of the DHT algorithms were invented in the early 2000s, with Kademlia being one of them. DHTs mostly differ just by how distance is defined, and how the neighbors are chosen. [15] DHTs have been developed to remove bottlenecks in peer search.

### 3.2.1 Kademlia

Kademlia is a DHT designed by Petar Maymounkov and David Mazières in 2002. It is based on a tree of identifiers which are split across peers on a network. The identifiers are 160 bits, e.g. a SHA-1 hash of some larger data. Kademlia tries to improve upon previous DHT-based routing algorithms by introducing a symmetric XOR metric for distance between node IDs in the key space. [16] These IDs are sorted in a binary search tree, with each node's position determined by the shortest unique prefix of its ID, like shown in diagram 3.2 on page 9. Kademlia makes sure that any node in the network can locate any other node by its ID by making sure that each node knows at least one of the nodes in each subtree.

A single query in Kademlia has been shown in real-world tests to result in an average of 3 network hops, meaning that the query gets relayed through two peers before reaching the requested resource. [17] Network hops are a necessary evil in distributed systems,

and Kademlia does well in requiring on average a  $\log(n)$  queries in a network of  $n$  nodes. Since the closeness metric is based on a similarity search rather than a measurement, the closest peer is only closest by the identifier, not by network latency. [18]

The randomness of Kademlia is great at averaging the network hops required to reach a scarce resource. The downside is that it also averages everything else, increasing latency to closest connected peers, and increasing the minimum hops to reach a common resource.

Kademlia protocol has four remote procedure calls, or RPCs in short. These are PING, STORE, FIND\_NODE, and FIND\_VALUE. A Kademlia participant's most important operation is node lookup, which is locating  $k$  closest nodes to a given node ID. It is a recursive operation, which starts by picking  $\alpha$  closest nodes from its closest non-empty bucket, and sending them all FIND\_NODE calls. This is repeated until the initiator has queried and got responses from all  $k$  closest nodes it has seen.

### **Eclipse Attacks**

Although Kademlia is usually used with random hashed identifiers, the distance metric stays the same. By forging identities that are close-by, you can advertise false friends which take over the search space. For example, in a 2019 paper "Eclipsing Ethereum Peers with False Friends" by S. Henningsen, D. Teunis, M. Florian, and B. Scheuermann, they demonstrate that to eclipse a victim in Ethereum P2P network, you need to fill its 8 slots for outbound connections, and fill 17 slots for inbound connections to completely deny service without going through the attacker's nodes. [19]

## Chapter 4

# Verifiable Delay Functions

A verifiable delay function is a function that calculates sequential calculations that cannot be significantly sped up by parallelisation or skipped, and creates a proof of the calculations that is faster to verify than the calculations are to fully re-evaluate. The computation of a VDF can be split into three distinct parts: Evaluation, proof generation, and verification. The evaluation is the hard sequential calculation from which a proof is generated.

Verifiable delay functions are a peculiar bunch of cryptographic proofs, since they are not proving that something has happened or something has been seen, but that a certain amount of time has been spent to compute the result. This results in the fact that in many use cases the starting parameters for VDF calculation must be given by a third party or computed by a multi-party computation. This is to prevent possible precomputation — meaning that the peer that provides a VDF proof could have computed it earlier, separately from the current transaction.

There's no known way to reduce computation costs sublinearly to the bit length of the exponent used. [20] This means that there's no algorithm-based solution to make VDF computation faster that is not related to other factors like efficiency. Despite this, there is a growing effort in developing faster hardware for VDF evaluation, especially for doing repeated modular squarings.

## 4.1 History

Verifiable delay functions are based on time-lock puzzles. Like VDF's, time-lock puzzles are computational sequential puzzles that require a certain amount of time to solve.[21] Time-lock puzzles were originally created to send information into the future, which simplifies into a cryptographic time capsule. Ronald L. Rivest, Adi Shamir, and David A. Wagner classify time-lock puzzles generally under the term timed-release crypto in their 1996 paper. [21] They describe the envisioned use cases for timed-release crypto as being the following, and they're very similar with the general concept of sending information into the future:

- A bidder in an auction wants to seal his bid so that it can only be opened after the bidding period is closed
- A homeowner wants to give his mortgage holder a series of encrypted mortgage payments. These might be encrypted digital cash with different decryption dates so that one payment becomes decryptable and thus usable by the bank at the beginning of each successive month
- An individual wants to encrypt his diaries so that they are only decryptable after fifty years.
- A key escrow scheme can be based on timed-release crypto so that the government can get the message keys but only after a fixed period, say one year

Verifiable delay functions fall under the same definition, but introduce a publicly verifiable proof that is much faster to verify than the puzzle was to solve. The original motivation for verifiable delay functions was to enable someone to pick off where somebody else had left when calculating a time puzzle — to make it possible to create checkpoints in the calculation that could be trusted. At the same time, VDFs found use in blockchain

applications as a way to achieve unspoofable randomness in elections and smart contract input.

## 4.2 Use Cases

Many have shown that there are more use cases for these algorithms than puzzles and random number generators. Some examples include preventing front running in P2P cryptocurrency exchanges, spam prevention and rate limiting [22] Almost all use cases are based on the property that with a new unique input, and a difficulty requirement, there is no easy way to speed up the calculation.

In the front-running use case, this property can be used to restrict front-running in both centralized and decentralized exchanges. [23] Front-running is a term used in both stock trading and cryptocurrency trading which means using inside knowledge of a future transaction to trade the asset with a better price or deny the other transaction from happening. [24] The issue is pronounced in trustless P2P networks like the Ethereum blockchain, where pending transactions are for all to see and can be overtook by giving the network validators a bigger reward, for example. [25] In a centralized exchange, preventing front-running with a VDF means that the exchange waits for a specified time before it starts the fulfilling the trade. The VDF serves as a receipt for the order taker that there was no reordering and the take orders were processed in a FIFO<sup>1</sup> fashion. [26] In a decentralized exchange the same holds, if there exists only one peer that can fulfill the order. The peer then functions exactly like the centralized exchange mentioned before.

Spam prevention and rate limiting in the context of VDFs mean roughly the same thing. This is because spam prevention with VDFs is done with rate limiting. If a network application requires peers to perform some sort of Proof-of-Work in a form of a VDF, by defining the difficulty parameter  $T$  they can set a time boundary for subsequent

---

<sup>1</sup>First in, first out

requests, also called a rate limit. If all the peers need to perform a challenge to request resources from the requestee, a performance-based limit is created for the time between each individual request, thus limiting spam.

### **4.3 Variations**

There are multiple variations on the security principles used, since a VDF can use any group of unknown order as a basis for its unpredictability. Known solutions use RSA, elliptic curves and class groups as their cryptographic basis.

## **4.4 Similar Constructs**

### **4.4.1 Proto-VDF**

A VDF can only be calculated sequentially, but even without a proof there is a possibility to make the verification faster through parallelism. A non-verifiable delay function, or time-lock puzzle in short, can be still verified faster than the calculation, because there is no sequential requirement after the puzzle has been calculated, enabling to use multiple CPU cores or highly parallel graphics processing units for verifying the puzzle, like in Solana. [6] One could also argue that all blockchains and some random beacons form multi-party calculated VDFs. Solana's Proof of History can be described as a proto-VDF, since it was invented before the term VDF was coined, but with similar guarantees.

### **4.4.2 Classic Slow Functions**

Since VDF's idea is that a hard calculation can be verified faster than to be evaluated, any calculation of which inverse is harder is a candidate for a VDF-type construct. One example of these asymmetric calculations is Sloth. [20] Sloth doesn't include a proof

and is not asymptotically verifiable, thus not filling the definition of a VDF per se.<sup>2</sup> The evaluation of Sloth is calculated with repeated cubic roots, and is verified with repeated cubings, with cubic roots being harder to compute than the cubings.

## 4.5 Hardware Developments

VDF applications can be made faster with hardware, and it has been estimated that with an ASIC<sup>3</sup> chip a VDF can be evaluated more than ten times faster than with a GPU. [27] The three different parts of a VDF calculation are different in terms of the hardware that can be utilized to make them faster. The evaluation part currently has no de facto hardware to make it faster, but the ASIC and FPGA developments are the ones pushing this part forward. Proving can be made faster with GPUs, whilst verification benefits from fast general computation, thus CPUs.[28] The chip developments for evaluation also help calculation of zero-knowledge proofs, since they also use repeated squarings for proof generation, thus benefiting cryptography at large. If or when hardware specifically optimized for sequential squarings is commercialized, VDFs can become much more mainstream, and suffer less from computational differences, thus requiring less trust between the calculating parties. It's a race against time, since if an attacker put a considerable amount of resources in the development of such a chip themselves, they could potentially compromise the security of some blockchains, due to an assumption called weak subjectivity. [29] Weak subjectivity means favoring a blockchain with the most blocks in it, rather than focusing on a majority vote. [30] By creating blocks faster than the majority of the network, one could attack proof of stake systems relying on weak subjectivity with a proprietary chip.

Besides the development of an ASIC driven by vdfresearch.org [22], there's been de-

---

<sup>2</sup>Asymptotic, meaning that there is a maximum verification time and the verification can't scale beyond that point.

<sup>3</sup>Application-Specific Integrated Circuit. A purpose-built chip for a single algorithm.



velopment on CPU instructions by Intel, aiming to help the specific operation of modular exponentiation, which is used not only in VDFs, but also in classical RSA, DSA, and DH algorithms, not to mention homomorphic encryption. [31]

# Chapter 5

## Proof of Latency

Proof of Latency, "the algorithm" or "PoL", is an algorithm that when used in a P2P context, can offer a robust way of reducing network latency between peers by minimizing the number of hops between peers that are close in terms of performance and geographical location. PoL also makes network bootstrapping faster, making it easier to find performant, close-by peers on first interaction with PoL-enabled network. PoL can result in a network in which a peer can estimate its latency to another peer before connecting to it.

The algorithm is trustless and requires no specific hardware from the participants, while a trusted computing platform would make it fairer and more reliable by not discriminating based on processing performance. An ASIC VDF chip would also make any of the later mentioned attack vectors less relevant by removing variance in hardware.

### 5.1 Use Cases

The section Peer-to-peer Networking described the problems Proof of Latency is trying to solve. While the use in P2P routing might be a no-brainer, the same algorithm can be used for computation performance benchmarking. In the next sections I will describe these use cases more thoroughly.

### 5.1.1 Dynamic P2P Routing

The use case that drove me to proceed working on Proof of Latency is an idea that there could be a trustless way of telling to other peers how much latency you have to another peer. In current P2P systems, if one peer were to tell you that it had a 10ms latency to another peer, there would be no guarantees of the 10ms latency holding true. If we could make a proof of that latency using cryptography, you could tell that the reported latency is true, and some work has been done to calculate it. This kind of proof could also make eclipse attacks harder to accomplish by requiring a closer distance or more resources to reach a lower latency to the targeted peer.

Although DHTs like Kademlia do peer distribution basically at random based on identifier closeness, there are no guarantees that when a peer connects to peers it has received from an another peer are also random, and thus the promise of random peer walk is lost. With Proof of Latency, you can be sure that the peers advertised are actually close-by. Of course, this is easily avoided by the attacker by spawning multiple peers in the same location.

### 5.1.2 Benchmark Queries

A protocol like this could be used to query for performance. In fact, Proof of Latency does that already, but processor development might render that property less effective in the future, unless one were to measure parallel processing capability with multiple VDFs, for example. Parallel multiple VDFs have been thought of and tested before, but calculating multiple separate VDFs has not been useful in previously imagined use cases, since it doesn't serve as a measure of time, but performance.

This could serve as a part of a greater protocol in a distributed computation system. It would be very unfortunate to run large datasets against a data analysis model on a mobile phone, and it could be beneficial to prove that the peer that advertises its services can actually run the computation without hiccups. There have been proposals for a VDF-as-a-

function system, in which less performant peers could query for a VDF calculation if they don't have the means to do so in a time frame small enough themselves. The FPGA based system is being tested right now on Amazon Web Services cloud platform. {Devlin2020-qw A benchmark query using PoL could also be a part of such a system to verify peers' ability to perform VDF calculations faster than the querying peer.

## 5.2 Role of Latency in Distributed Systems

It's hardly a surprise, but latency is a huge factor in distributed systems, especially trustless, decentralized ones. Latency is mostly constrained by the speed of light, which can not be changed, and thus there are concrete factors that must be taken into account when designing a P2P system with routing. In 2012, the global average round-trip delay time to Google's servers was around 100ms. [32]

In the new space age the maximum possible latency grows very fast, as there could be peers joining to a distributed network from other planets, space ships or stations. This might be unnecessary to think about in the distributed P2P context for now, but before all that, we have global satellite mesh internet providers, like Starlink. Elon Musk, the founder of SpaceX, which provides the network of satellites, claims that there's going to be a round-trip<sup>1</sup> latency of about 20 milliseconds between a single satellite and the user. [33] In legacy satellite internet access, the round-trip time even in perfect conditions is about 550 milliseconds. [34] This difference between legacy and newer satellite internet comes from the difference in their orbits and the sheer amount of satellites involved. Legacy satellite internet uses geostationary orbits, which are very high, beaming on a single face of the earth at a time with limited bandwidth. Newer systems, like Starlink, use a low-earth-orbit, which requires more satellites, since they zoom by at such a speed that constantly changing which satellite you're connected to is a must. The low orbit

---

<sup>1</sup>Including the user's initial request and received response

also means less distance between the satellites and the user. The 20 millisecond latency claimed by Starlink at first seems like a stretch, but is believable when you take into account that inter-satellite links are done by laser, and light can travel about 31 percent faster in a vacuum than in fiber optics. [35] Intercontinental latencies can become much lower because of this. In blockchains, latency plays a role in the efficiency of the power used to achieve consensus. Miners waste energy on a previous block as long as they don't receive information on the winner of the previous block race. Simulations by Wei Bi, Huawei Yang, and Maolin Zheng in their paper *An Accelerated Method for Message Propagation in Blockchain Networks* have shown that if you calculate the round-trip time between the peers that are connected to each other and dropping the ones with larger latencies in favor of lower ones, you can achieve 50% improvement in average latency with 1 to 2 peers connected. When connectedness grows from the degree of just 1-2 peers up to 20 connected peers, the average latency improvements achieved drop to about 20%. [36] You can't keep multiplexing connections<sup>2</sup> forever, though, and there's a Goldilocks zone for the most effective amount of connections. When connectedness increases, there's shorter routes simply by chance to peers you're not directly connected to, and protocols like publish-subscribe schemes work faster, propagating their messages to the whole network faster because there's less relaying happening. There are hardware and software related limitations to the amount of peers. On IPFS, for example, the protocol has been breaking user's routers [37] because of the high number of incoming connections that need to be routed through NAT<sup>3</sup>.

---

<sup>2</sup>Having multiple concurrent stateful connections.

<sup>3</sup>Network address translation. Hides the local area network from the internet under a different subnet address.

### 5.3 Network Hops Increase Latency

Network hops in P2P systems are introduced when two peers are not directly connected to each other, but rather through one or many relays. There are network hops that cannot be easily avoided, like the hops between network routers in the internet. Most of the P2P routing protocols used today are oblivious to the problem of introducing large hops to communications between two peers, trading network performance for network robustness and decentralization. These DHT-based protocols, like Kademlia, make the assumption that their users have fast internet access, and minimize the average latency by selecting connected peers basically at random.

While the randomness is great for preventing eclipse attacks, they can introduce unnecessary geographical hops between two peers. If two peers are in the same WAN, for example, in Kademlia they might still connect to each other through a network hop going through another continent. This makes individual connections less efficient. Now, if we were to rely on IP address geolocation, we could more efficiently connect to peers that are close-by. This is unfortunately impossible in privacy-oriented P2P networks, like mixnets, which aim to hide as much of the packet routing information as possible, by routing individual packets through different peers and hiding IP addresses of two connected peers from each other. [38]

Proof of Latency is made to improve the performance of current P2P networking solutions and make them future-proof, even for hops between planets, while still being compatible with some privacy-preserving P2P protocols, since it is agnostic of the addressing method used.

### 5.4 Protocol Description

Proof of Latency is an interactive public-coin protocol that produces a publicly verifiable non-interactive proof by publishing the parameters with the proof, signed by the partici-




Figure 5.1: Protocol Diagram, Proof of Latency

pants. The protocol cannot be made non-interactive because of the requirements involved with The most important part when calculating any VDF is the setup. If a prover wants to create a valid proof, it needs a previously unknown starting point at first, so that it can prove for certain that it cannot have calculated the VDF in advance. Proof of Latency tackles this problem with a two-way RSA Diffie-Hellman key setup with random ephemeral prime numbers, which is an industry standard to achieve a previously unknown key with large certainty.

The proof generation removes trust between the two peers by introducing a race between them. The two participants are called the prover and the verifier, although they both are actively proving and verifying their VDFs. The difference here is that they both calculate a VDF and the verifier calculates the difference in iterations between the two.

First the prover and the verifier do a Diffie-Hellman key exchange to construct a previously unknown key. Then, they both start calculating a VDF in parallel. The prover only calculates the VDF up to a predefined threshold, and then sends the proof with a signature back to the verifier. The verifier then stops its own calculation, generates a proof of it, and then calculates the absolute difference between the amount of iterations between its own VDF and the prover's.

Since calculating a VDF is relatively easy for modern processors, a VDF over as little as a few milliseconds of time can be a valid way of measuring latency. Still, without an ASIC chip for calculating VDFs faster than any other available processor, these protocols are also a measurement of processing performance. This might introduce an unfortunate barrier for entry for mobile and IoT devices. The second version of PoL is meant to tackle this problem by creating a performance and latency gradient to the network. The network topology results in a gradient that is defined by geographical location and the similarity in performance. This means that connectedness between mobile and IoT devices is going to be better than between devices that have a huge performance difference.



## 5.5 Attack Vectors

Since Proof of Latency removes security guarantees by removing randomness from routing, some new attack vectors are introduced. The two versions of the algorithm differ much in terms of security. Proof of Latency is not meant to be a one-stop shop towards a safer network, but rather an add-on to make things more efficient while still keeping security in mind. The following attack vectors are a product of prior knowledge, research and the writer's thought process, and none have been tested out in practice, yet.

### 5.5.1 Advertising Dishonest Peers

A sybil network participant can spawn an arbitrarily large number of new identities and network peers that are close-by, create multiple Proofs of Latency with them, and only advertise these peers to the rest of the network on their DHT. Two peers can also fake the proof if they co-operate by using VDFs they have already calculated and matched their outputs to be as close as possible.

This can be mitigated somewhat by requiring peers to renew their proofs regularly or by using a public source of randomness. It's hard to make systems completely trustless, and this was a good example while doing this research. Unless the proofs are based on a blockchain, there's no "blind" trust that can be achieved with PoL on how low a latency another peer has to another unless it is calculated yourself by a process you control. Even in this case, the proofs would need constant updating and would at most have a minimum trusted latency roughly equal to the blocktime. This doesn't leave many chains out there at the time of writing that have the capability to do this, and even them would maybe have to introduce some localized consensus to decrease latency even more.

While this sounds bad, if the protocol were to be enforced in a way that it required a proof to be generated between an advertiser and every advertised peer, these sybil, dishonest peers would not propagate falsely to the rest of the network but would remain a

separate blob.

## 5.5.2 Performance Matching

Both versions suffer from the same issue, let's call it performance matching, which is a timing attack. Timing attacks are a family of attacks against computer systems called side-channel attacks, which attack the fact that although software is abstract and can be very well fit against any attacks on the software level, outer factors still affect the hardware it runs on. Timing attacks rely on gathering of timing data from the target. [39] In performance matching, this means connecting to the targeted peer by another protocol or comparing existing proofs of latencies from all peers that have calculated their latency with the target.

Performance matching enables attackers to perform an eclipse attack on low-performance devices by matching the attacker's performance with the targeted mobile device so that it is as close as possible in the difference between iterations in PoL. Now again, if the algorithm had a trusted platform module requirement, or we had an ASIC for repeated squarings, this wouldn't be an issue. This attack could result in a complete network split.

## 5.6 Protecting Against Performance Matching

### 5.6.1 Hiding Information of VDF Results

Zero-knowledge proofs could be used to protect against performance matching. If the publicized proof didn't include both the VDF results and iterations, but just included the iteration difference, an attacker would have less info on each peer. This would make attacking more difficult, requiring more queries and PoL runs on average before finding a vulnerable peer.

### 5.6.2 Web of Trust

There's also a possibility of introducing a web of trust in parallel to PoL to recognize and shut out malicious peers more effectively. An example of such a system is SybilLimit, which adds a construction called trusted routes to DHT-based routing. [40] I see that trusted routes in conjunction with Proof of Latency would be a great couple. The problem of advertising dishonest peers could also be tackled with a trust-based system. Introducing too much trust into the system could also make PoL useless, since there could be a simple ping protocol with a time-based latency instead.

### 5.6.3 Peer Scoring

Peer scoring is used regularly in P2P networks, Proof of Latency serves as a peer scoring metric by itself, and can be used in various ways. To hinder any attacks, peers with the lowest latencies would be kept for a longer time even in the case of the peer not being online, and favoring new ephemeral connections that are farther.

### 5.6.4 Switch Responsibilities

This remedy aims to improve PoL version 1 by flipping the responsibilities of the prover and the verifier. The calculating party wants to prove its latency to the verifier by actually calculating the VDF itself, requiring the verifier to respond as fast as it can to the prover's messages. This way, only a proof of one VDF is needed, which the prover then advertises to the network, telling connecting peers exactly the number of calculations it got to calculate before the verifier responded to it.

There's a catch, though. The prover can still lie to other peers by slowing down their VDF calculation. By flipping the responsibilities the protocol becomes trustless for the prover, but not for others. The flip felt promising at first glance, but I soon realized there's no better way to do this than the second version, where trust is shared and not assumed.

### 5.6.5 Results

Since the network peers are trying to compete with each other for the least latency, they win the game by calculating the least amount of VDF iterations. In the first version of the algorithm, this soon becomes a problem. There's no way to tell how fast another computer is, so in this case spoofed results of 1 iteration would win every time.

This version only works if the peers can trust each other, but in that case we could also resort to using ping. Ways of introducing trust would be using a trusted computing platform, but even that could be vulnerable to undervolting or other physical tampering. This problem drove me to revert back to my original idea, which I'm calling the version two of Proof of Latency.

## Chapter 6

# Proof of Concept

To test out Proof of Latency, I made a software proof of concept in Rust. I picked Rust as a programming language because I've wanted to try it out in my own projects for long, and it has some good properties when it comes to cryptography and distributed computing. Since distributed computing or any server program for that matter uses some sort of concurrency to get non-blocking responses to requests they receive, you have a very good chance of running into race conditions with most systems programming languages. When a Rust program has been written in the default "safe" compiler setting, a data race where two or more threads try to access a shared memory resource at the same time is simply impossible to achieve. [41] Also, as it is a compiled language without a garbage collector, and with lots of optimization in the compiler, Rust has good baseline performance even when compared to other similarly low level programming languages like Go. [42] Rust has seen a surge of interest in the last few years, and is used in many projects in production, notably in embedded and distributed computing, the former because of Rust's modern build tooling and robustness.

William Borgeaud's blog post [43] from November 2019 in particular was the first reference I encountered that described verifiable delay functions in familiar terms to a software engineer like me. The blog post and its accompanying code [43] that also happened to be in Rust helped me bootstrap the project.

First of all, I knew I had to make the VDF run asynchronously, or if possible, in another thread, because it would be one of two VDFs needed in the protocol. There needs to be a separate task for listening to the other party's input, and to make it possible to test locally without networking. For the VDF that is used in Proof of Latency it was also needed that the prover's calculation could be ran indefinitely and ended abruptly by a received "cap" prime number from the verifier. No previous VDF library seemed to have this option, because the difficulty parameter  $T$  is usually predetermined in contrast to Proof of Latency.

The code is structured as follows: A runnable binary demo (main.rs) and the reusable library it also uses (lib.rs) resides in the top crate. They both depend on the subcrates "P2P" and "vdf". Here is an example output of Proof of Latency run:

---

```

pub fn run_vdf_worker(
    self,
) -> (Sender<Int>, Receiver<Result<VDFProof, InvalidCapError>> >) {
    let (caller_sender, worker_receiver): (Sender<Int>, Receiver<Int>) =
        channel();
    let (worker_sender, caller_receiver) = channel();

    thread::spawn(move || {
        let mut result = self.base.clone();
        let two = Int::from(2);
        let mut iterations: u32 = 0;
        loop {
            result = iter_vdf(result, &self.modulus, &two);
            iterations += 1;

            if iterations == self.upper_bound || iterations == u32::MAX {
                // Upper bound reached, stops iteration and calculates the proof
                debug!(
                    "Upper bound of {:?} reached, generating proof.",
                    iterations
                );

                // Copy pregenerated cap
                let mut self_cap: Int = self.cap.clone();

                // Check if default, check for primality if else
                if self_cap == 0 {
                    self_cap = Generator::new_safe_prime(128);
                    debug!("Cap generated: {:?}", self_cap);
                } else if !self.validate_cap(&self_cap, iterations) {
                    if worker_sender.send(Err(InvalidCapError)).is_err() {
                        error!("Cap not correct!");
                    }
                    break;
                }

                // Generate the VDF proof
                let vdf_result = VDFResult { result, iterations };
                let proof = VDFProof::new(
                    &self.modulus,
                    &self.base,
                    &vdf_result,
                    &self_cap,
                );
                debug!("Proof generated! {:?}", proof);

                // Send proof to caller
                if worker_sender.send(Ok(proof)).is_err() {
                    error!("Failed to send the proof to caller!");
                }
            }
        }
    });
}

```

To achieve a better algorithm design, the library itself is described as a state machine, which also helps in thinking of the algorithm in TLA+. Variable names have been converted into more readable non-mathematical forms.

The following is a code example of the threaded evaluation part of a VDF inside proof of latency, containing evaluation stop logic for both the prover and the verifier:



---

```

pub fn run_vdf_worker(
    self,
) -> (Sender<Int>, Receiver<Result<VDFProof, InvalidCapError>>) {
    let (caller_sender, worker_receiver): (Sender<Int>, Receiver<Int>) =
        channel();
    let (worker_sender, caller_receiver) = channel();

    thread::spawn(move || {
        let mut result = self.base.clone();
        let two = Int::from(2);
        let mut iterations: u32 = 0;
        loop {
            result = iter_vdf(result, &self.modulus, &two);
            iterations += 1;

            if iterations == self.upper_bound || iterations == u32::MAX {
                // Upper bound reached, stops iteration and calculates the proof
                debug!(
                    "Upper bound of {:?} reached, generating proof.",
                    iterations
                );

                // Copy pregenerated cap
                let mut self_cap: Int = self.cap.clone();

                // Check if default, check for primality if else
                if self_cap == 0 {
                    self_cap = Generator::new_safe_prime(128);
                    debug!("Cap generated: {:?}", self_cap);
                } else if !self.validate_cap(&self_cap, iterations) {
                    if worker_sender.send(Err(InvalidCapError)).is_err() {
                        error!("Cap not correct!");
                    }
                    break;
                }

                // Generate the VDF proof
                let vdf_result = VDFResult { result, iterations };
                let proof = VDFProof::new(
                    &self.modulus,
                    &self.base,
                    &vdf_result,
                    &self_cap,
                );
                debug!("Proof generated! {:?}", proof);

                // Send proof to caller
                if worker_sender.send(Ok(proof)).is_err() {
                    error!("Failed to send the proof to caller!");
                }
            }
        }
    });
}

```

The following is a code example of the proving part of a single VDF:

---

```
pub fn new(  
    modulus: &Int,  
    generator: &Int,  
    result: &VDFResult,  
    cap: &Int,  
) -> Self {  
    let mut proof = Int::one();  
    let mut r = Int::one();  
    let mut b: Int;  
    let two: &Int = &Int::from(2);  
  
    for _ in 0..result.iterations {  
        b = two * &r / cap;  
        r = (two * &r) % cap;  
        proof = proof.pow_mod(two, modulus) * generator.pow_mod(&b, modulus);  
        proof %= modulus;  
    }  
  
    VDFProof {  
        modulus: modulus.clone(),  
        generator: generator.clone(),  
        output: result.clone(),  
        cap: cap.clone(),  
        proof,  
    }  
}
```

---

The following is a code example of the verification of a VDF proof:

---

```
pub fn verify(&self) -> bool {  
    // Check first that the result isn't larger than the RSA modulus  
    if self.proof > self.modulus {  
        return false;  
    }  
    let r =  
        Int::from(2).pow_mod(&Int::from(self.output.iterations), &self.cap);  
    self.output.result  
        == (self.proof.pow_mod(&self.cap, &self.modulus)  
            * self.generator.pow_mod(&r, &self.modulus))  
            % &self.modulus  
}
```

---

The proof of concept is made as a library, so that it can be imported in other projects and tinkered with. To make it into a library in addition to a stand-alone binary, I used structs for logical instances of a single VDF, for example. This is mainly to keep program state in instances of these structs, removing possible side effects.

## 6.1 Tests

### 6.1.1 Unit Tests

To make sure that the VDF calculations worked on arbitrary input by generating a successfully verifiable proof I made unit tests for the VDF evaluation, proof generation and verification. To further ensure the soundness of the algorithm, I made simulated cases for proof of latency with unit tests with separate, async threads for each execution. Unit tests reside in the same file the program code itself, and the Rust standard library has well enough functionality so that I wasn't forced to use a separate library for assertions<sup>1</sup>, for example.

---

<sup>1</sup>Checking that a statement holds true or is equal to something

In addition to regular unit tests with predefined input, I dabbled in property testing. Property testing is a term that belongs somewhere between fuzzing<sup>2</sup> and unit tests. Otherwise it functions like regular unit tests, but it adds generated input into the equation. By adding generated, sometimes randomized input, one can check more thoroughly for edge cases. While being effective in recognizing some unexpected bugs, it still has a way to go when compared to formal verification, which is a testing method that aims to define a system with mathematic proofs, verifying that the output should always, apart from hardware issues, be predictable and deterministic.

### 6.1.2 End to End Tests

To test the whole demo out in a simulated P2P setting I used Protocol Labs' Testground software. It is a tool to simulate a network with thousands of peers on one machine.

---

<sup>2</sup>Testing systems against highly randomized and a high volume of input.

# Chapter 7

## Future Considerations

If this system was integrated to a blockchain or a publicly verifiable source of randomness for the initial setup, the proofs could be verified by anyone against consensus. Not only this would add trust to the latency measurements, but also speed up initial bootstrapping of the P2P network. When P2P networks eventually grow larger and larger, the network bootstrapping infrastructure needs to be rethought to handle more traffic and be faster in its initialization. By getting introduced to the closest peers possible right at the start the user can experience a more performant network right from the beginning, lowering the barrier for entry by making first impressions better.

Also, I believe the work is not cryptographically as secure as it could be, and the field is progressing at a mindnumbing pace. Making sure the algorithm is VDF-agnostic would be a logical next step, since the field is still in progress of finding the best possible formulation of a VDF. Quantum computing can render all existing VDF types insufficient, and new VDF types could change the parameter logic fundamentally. The idea behind this thesis is to think of a new creative way of using verifiable delay functions by defining a protocol, and not to necessarily use the most robust cryptography available.

# Chapter 8

## Conclusion

Since calculating a VDF is relatively easy for modern processors, a VDF over as little as a few milliseconds of time can be a valid way of measuring latency. Still, without an ASIC chip for calculating VDFs faster than any other available processor, these protocols are also a measurement of processing performance. This might introduce an unfortunate barrier for entry for mobile and IoT devices. The second version of PoL is meant to tackle this problem by creating a performance and latency gradient to the network. The network topology results in a gradient that is defined by geographical location and the similarity in performance. This means that connectedness between mobile and IoT devices is going to be better than between devices that have a huge performance difference.

Proof of Latency as a peer scoring metric not only protects peers from eclipse attacks, but can also function as a way of speeding up the initial bootstrapping process by bringing peers more closely together.

# References

- [1] Gunnar Kreitz and Fredrik Niemelä. Spotify – large scale, low latency, P2P Music-on-Demand streaming.
- [2] Satoshi Nakamoto. Bitcoin: A peer-to-peer electronic cash system. Technical report, Manubot, 2019.
- [3] Benjamin Wesolowski. Efficient verifiable delay functions. Technical Report 623, École Polytechnique Fédérale de Lausanne, 2018.
- [4] Krzysztof Pietrzak. Simple verifiable delay functions. Technical Report 627, IST Austria, 2018.
- [5] Dan Boneh, Benedikt Bunz, and Ben Fisch. A survey of two verifiable delay functions. page 13.
- [6] Anatoly Yakovenko. Solana: A new architecture for a high performance blockchain, October 2018.
- [7] Jean-Philippe Aumasson. *Serious Cryptography*. William Pollock, 2018.
- [8] Stuart Cheshire and Marc Krochmal. Multicast dns. Technical report, RFC 6762, February, 2013.
- [9] pdp. Name (mDNS) poisoning attacks inside the LAN. <https://www.gnucitizen.org/blog/name-mdns-poisoning-attacks-inside-the-lan/>, January 2008. Accessed: 2020-6-3.

- 
- [10] GoTenna. Extend the edge of connectivity. <https://gotenna.com/>. Accessed: 2020-9-13.
- [11] Helium. Helium – introducing the people’s network. <https://www.helium.com/>. Accessed: 2020-9-13.
- [12] Mark Milian. Russians are organizing against putin using FireChat messaging app. *Bloomberg News*, December 2014.
- [13] Sam Biddle. The inventors of bluetooth say there could be problems using their tech for coronavirus contact tracing. <https://theintercept.com/2020/05/05/coronavirus-bluetooth-contact-tracing/>, May 2020. Accessed: 2020-9-24.
- [14] Protocol Labs. IPFS powers the distributed web. <https://ipfs.io/>. Accessed: 2020-9-24.
- [15] Xing Shi Cai and Luc Devroye. The analysis of kademlia for random IDs. *Internet Math.*, 11(6):572–587, November 2015.
- [16] David Mazières Petar Maymounkov. Kademlia. January 2020.
- [17] Stefanie Roos, Hani Salah, and Thorsten Strufe. Comprehending kademlia routing - a theoretical framework for the hop count distribution. *arXiv:1307.7000 [cs]*, July 2013.
- [18] Dean Eigenmann. From kademlia to discv5. <https://vac.dev/kademlia-to-discv5>, April 2020. Accessed: 2020-5-20.
- [19] Sebastian Henningsen, Daniel Teunis, Martin Florian, and Björn Scheuermann. Eclipsing ethereum peers with false friends. *arXiv:1908.10141 [cs]*, August 2019.
- [20] Dan Boneh, Joseph Bonneau, Benedikt Bünz, and Ben Fisch. Verifiable delay functions. In Hovav Shacham and Alexandra Boldyreva, editors, *Advances in Cryptology*



- *CRYPTO 2018*, Lecture Notes in Computer Science, pages 757–788, Cham, 2018. Springer International Publishing.
- [21] Ronald L Rivest, Adi Shamir, and David A Wagner. Time-lock puzzles and timed-release crypto. page 9.
- [22] VDF research. <https://vdfresearch.org/>. Accessed: 2020-2-11.
- [23] R Khalil, A Gervais, and G Felley. TEX-A securely scalable trustless exchange. *IACR Cryptology ePrint Archive*, 2019.
- [24] Dan Robinson. Ethereum is a dark forest - dan robinson - medium. <https://medium.com/@danrobinson/ethereum-is-a-dark-forest-ecc5f0505dfff>, August 2020. Accessed: 2020-9-21.
- [25] Cory Mitchell. What is Front-Running in stocks? <https://www.investopedia.com/terms/f/frontrunning.asp>, September 2020. Accessed: 2020-9-21.
- [26] Dan Cline, Thaddeus Dryja, and Neha Narula. ClockWork: An exchange protocol for proofs of non Front-Running. *MIT Digital Currency Initiative*, page 13, March 2020.
- [27] Stanford Video. Stanford BlockChain - day 2, 2020.
- [28] Protocol Labs, Kelly Olson. Hardware acceleration of RSA VDFs - kelly olson, September 2020.
- [29] VDF-based RNG with linear lookahead. <https://ethresear.ch/t/vdf-based-rng-with-linear-lookahead/2573>, July 2018. Accessed: 2020-2-19.
- [30] Vitalik Buterin. Ethereum wiki, weak subjectivity, August 2019.

- [31] Nir Drucker and Shay Gueron. Fast modular squaring with AVX512IFMA. In *16th International Conference on Information Technology-New Generations (ITNG 2019)*, pages 3–8. Springer International Publishing, 2019.
- [32] Ilya Grigorik. Latency: The new web performance bottleneck - igvita.com. <https://www.igvita.com/2012/07/19/latency-the-new-web-performance-bottleneck/>. Accessed: 2020-5-13.
- [33] Liam Tung. Elon musk: SpaceX’s internet from space should be good enough for online gaming. <https://www.zdnet.com/article/elon-musk-spacexs-internet-from-space-should-be-good-enough-for-online-gaming/>. Accessed: 2020-5-14.
- [34] Satellite internet latency - VSAT systems broadband satellite internet latency - broadband satellite internet service latency. <https://www.vsat-systems.com/satellite-internet-explained/latency.html>. Accessed: 2020-5-14.
- [35] Klint Finley, Will Bedingfield, Andy Greenberg, Medea Giordano, Brett Berk, Jess Grey, Matt Jancer, and Evan Ratliff. Brits approach (true) speed of light over fiber cable. *Wired*, March 2013.
- [36] Wei Bi, Huawei Yang, and Maolin Zheng. An accelerated method for message propagation in blockchain networks.
- [37] Whyrusleeping. Issue #3320, go-ipfs, “writeup of router kill issue”, October 2016.
- [38] Harry Halpin. nym-litepaper.pdf. <https://nymtech.net/nym-litepaper.pdf>. Accessed: 2020-5-14.
- [39] BearSSL - Constant-Time crypto. <https://www.bearssl.org/constanttime.html>. Accessed: 2020-5-18.

- 
- [40] Haifeng Yu, Phillip B Gibbons, Michael Kaminsky, and Feng Xiao. SybilLimit: A Near-Optimal social network defense against sybil attacks. In *2008 IEEE Symposium on Security and Privacy (sp 2008)*, pages 3–17, May 2008.
- [41] The Rust Project Developers. Races - the rustonomicon. <https://doc.rust-lang.org/nomicon/races.html>, November 2018. Accessed: 2020-9-24.
- [42] Jesse Howarth. Why discord is switching from go to rust - discord blog. <https://blog.discord.com/why-discord-is-switching-from-go-to-rust-a190bbca2b1f>, February 2020. Accessed: 2020-9-24.
- [43] William Borgeaud. understanding-vdfs. <https://wborgeaud.github.io/posts/understanding-vdfs.html>, November 2019. Accessed: 2020-3-9.